




KubeCon

North America 2024



CloudNativeCon



CLOUD NATIVE &
KUBERNETES

AI DAY

NORTH AMERICA

LLM Powered Agents with Kubernetes

Hema Veeradhi
Shrey Anand

Red Hat
Emerging Technologies (OCTO)



KubeCon



CloudNativeCon

North America 2024



Hema Veeradhi
Principal Data Scientist
Emerging Technology | Red Hat



California, United States



hemaveeradhi

Hi! 🖐️



Shrey Anand
Principal Data Scientist
Emerging Technology | Red Hat



Toronto, Canada



shrey-anand



Agenda

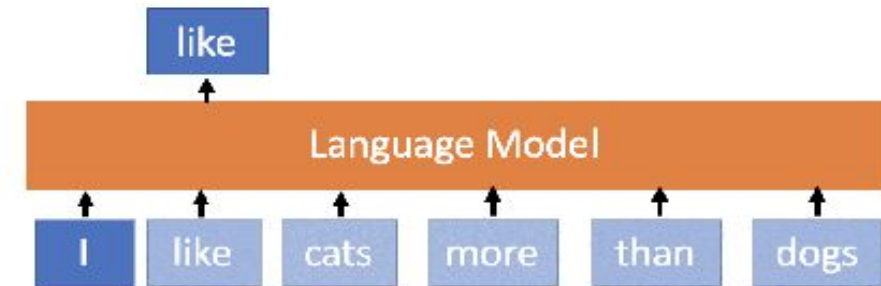
- **Introduction to LLMs**
- **Introduction to RAG**
- **What are Agents**
 - Components of LLM agents
- **Types of Agents**
 - Router based
 - ReAct
- **Agent Use Case**
- **Demo**
- **Concluding Remarks**

Introduction to LLMs

Large Language Models are **computational models that are capable of modeling and generating human language**. LLMs have the transformative ability to predict the likelihood of word sequences or generate new text based on a given input.

What are LLMs good at?

- Text generation/code generation
- Chatbots and Conversational AI
- Information retrieval
- Sentiment analysis



Input: n “tokens” -> I, like, cats, more, than

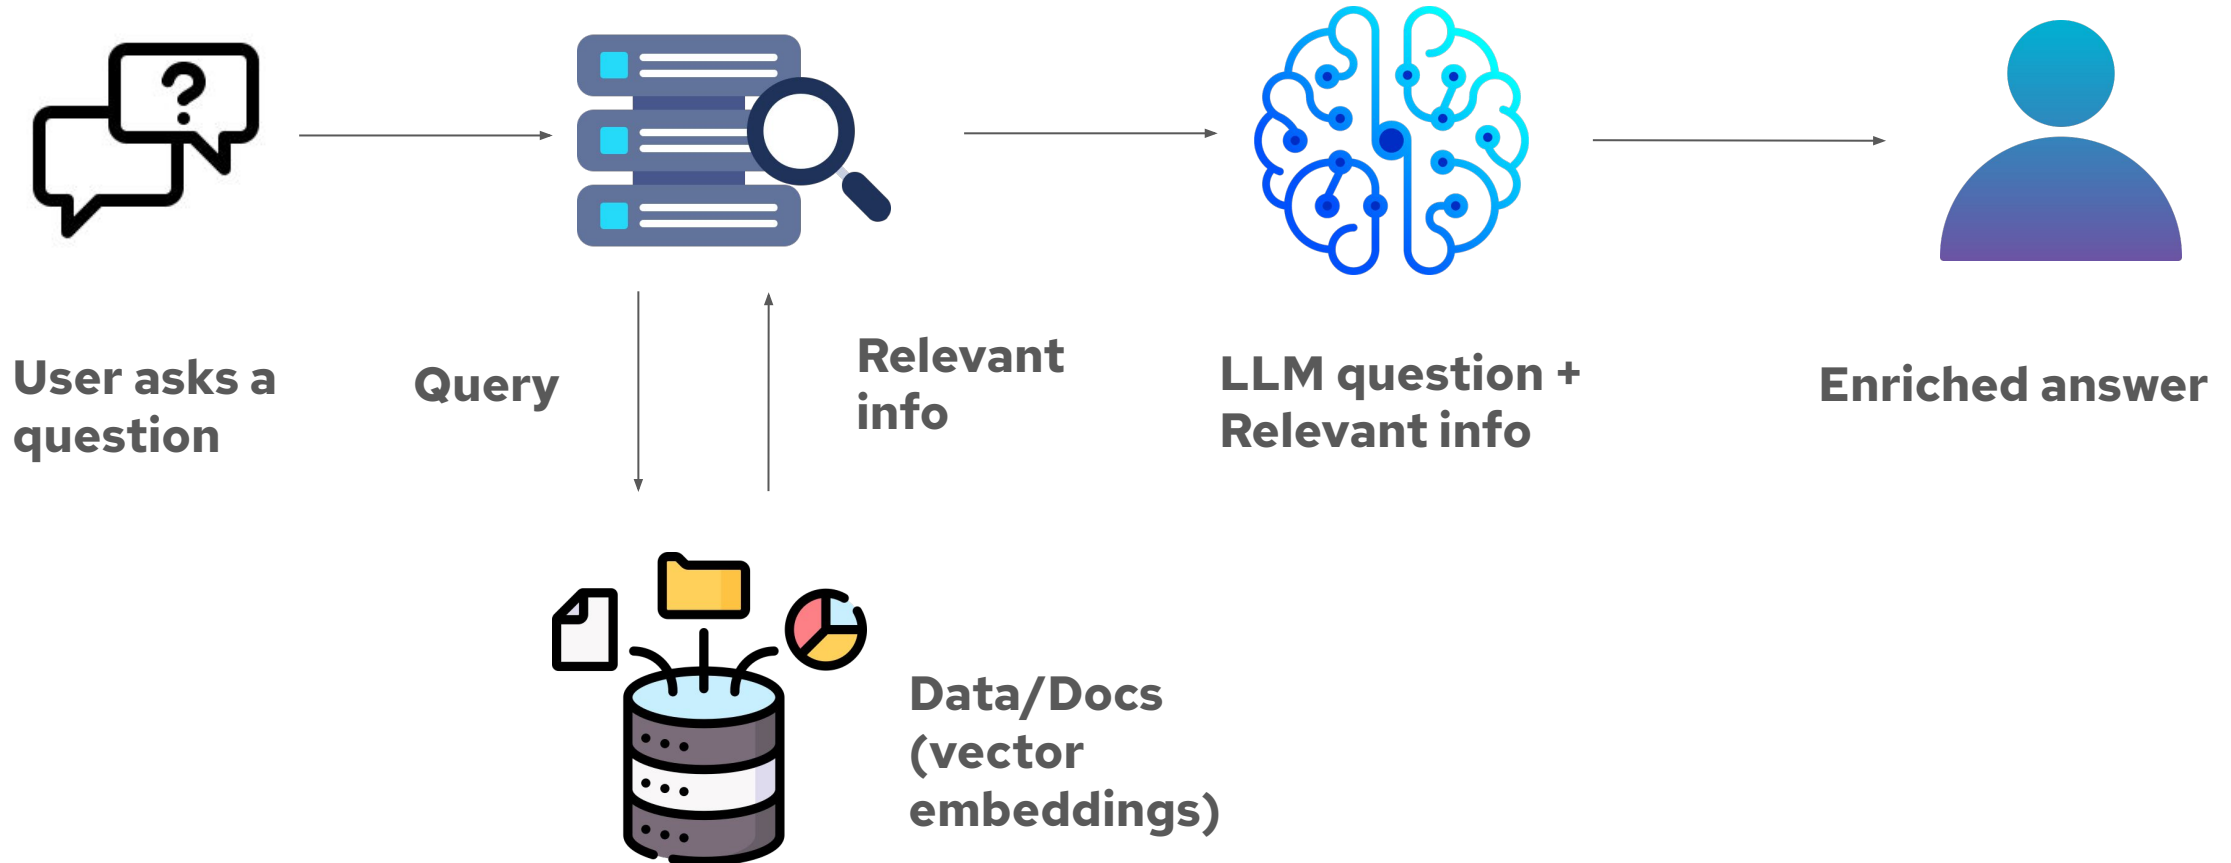
Output: 1 “token” -> dogs



When can LLMs fail?

- **Complex reasoning tasks** - LLMs have limited reasoning capability, LLMs are good knowledge retrievers but not good reasoners
- **No Dynamicity** - LLMs are static and unable to access real-time information
- **Limited Knowledge (hallucination)** - While trained on vast data, LLMs lack up-to-date world knowledge

RAG - Retrieval Augmented Generation





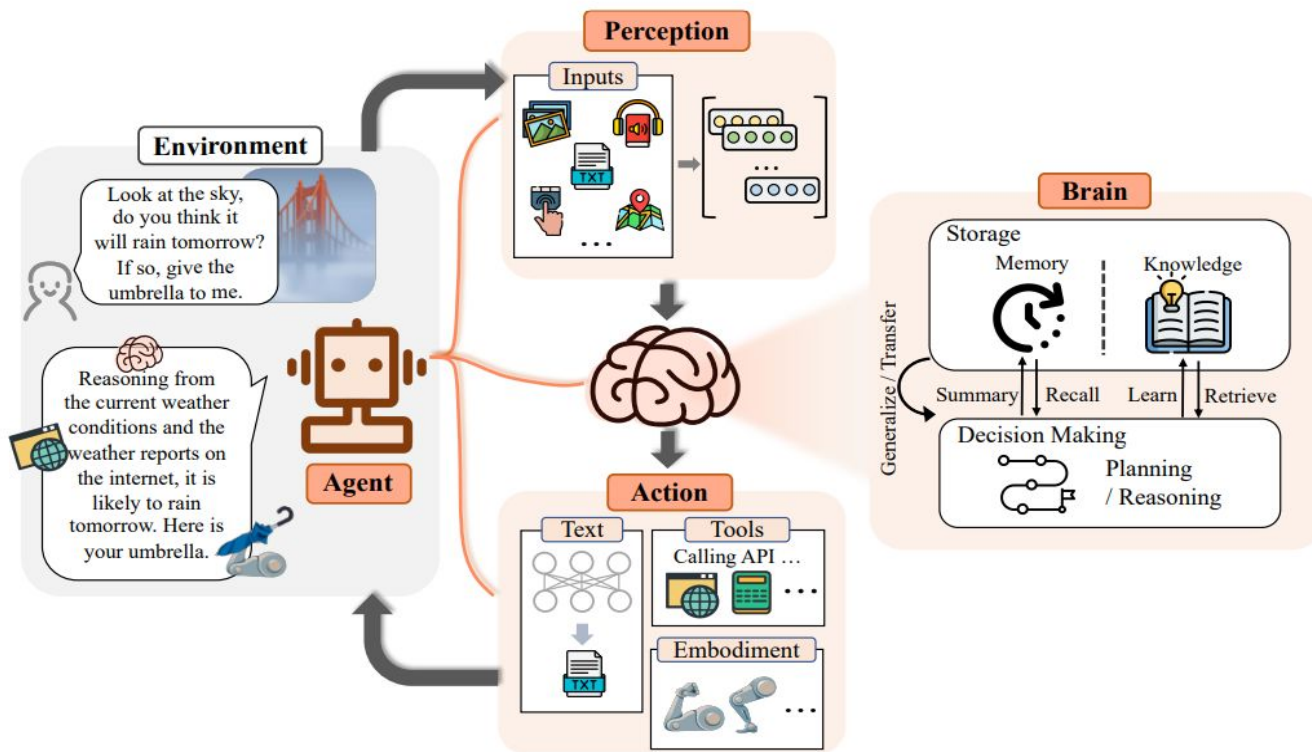
LLM Agents

LLM Agents, also known as Large Language Model agents, **leverage LLMs to execute complex tasks by integrating them with essential components like planning and memory.**

LLM Agent = LLM + Tools + State

- **LLM** - Computational engine i.e “brain”
- **Tools** - Agents ability to interact with the external world
- **Memory/State** - Agent’s memory of previous messages and results from used tools

Components of an LLM Agent



- **LLM** - Computational engine i.e. "brain"
- **Planning** - Chain of thought process (CoT) to create a plan for executing the tasks
- **Tools** - Executable functions, APIs, other services that complete various tasks
- **Memory** - Short term memory to retain the agent's train of thought and long term memory to retain context/conversation history
- **Actions** - Agents perform actions based on their environment and reasoning, adapting and solving tasks iteratively through feedback.



Types of Agents

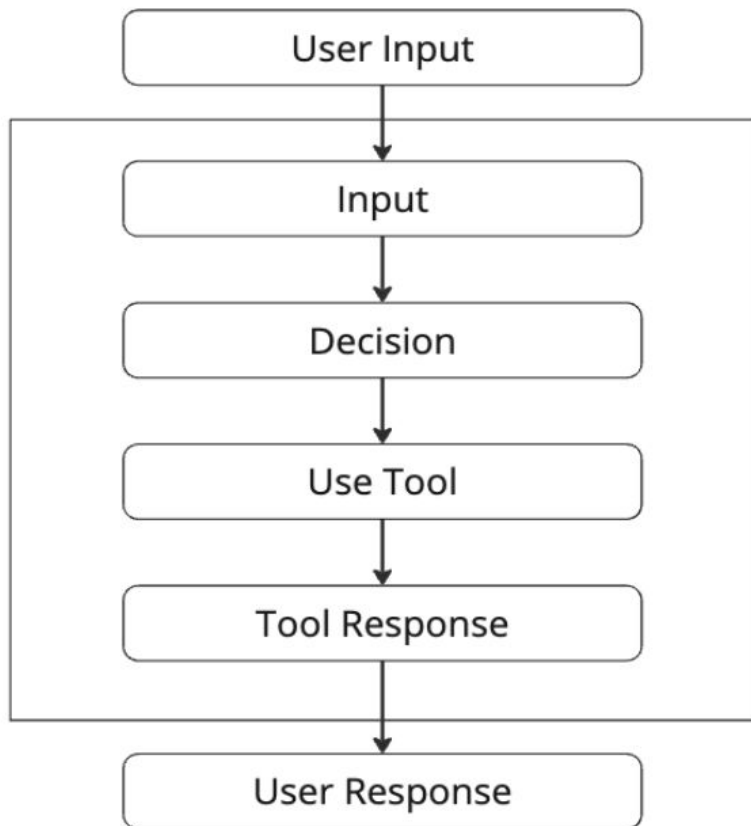
There are many different types of LLM agents to choose from, depending on the nature of your use case, including:

- React Agent
- Rule-based Agent
- Single/Multi Agent

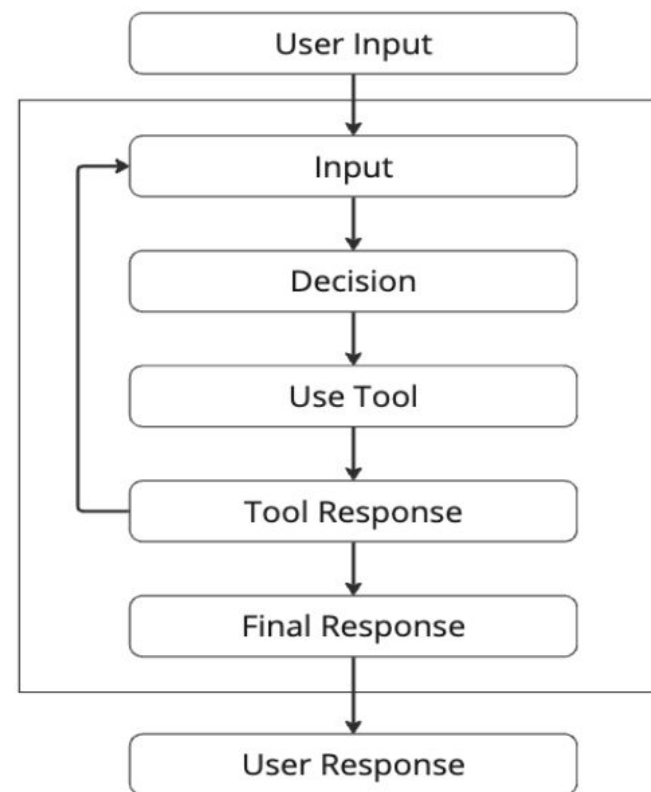
Tools and frameworks to build LLM agents

- [LangChain](#)
- [Bee Agent](#)
- [OpenAgents](#)
- [CrewAI](#)
- [Arch](#)

Router Agent



ReAct Agent



Use Case

- Imagine a **fictional cloud company CloudForge Dynamics that has several departments and is looking to have a unified chat assistant** for its employees to answer queries across these departments
- Let's say we have access to Products, HR, and Customer accounts documents
- How can we answer questions by gathering/combining information from all of these different sources?
 - A LLM agent can decide which department's sources should be used for answering the query




KubeCon



CloudNativeCon

North America 2024



**CLOUD NATIVE &
KUBERNETES**

AI DAY

NORTH AMERICA

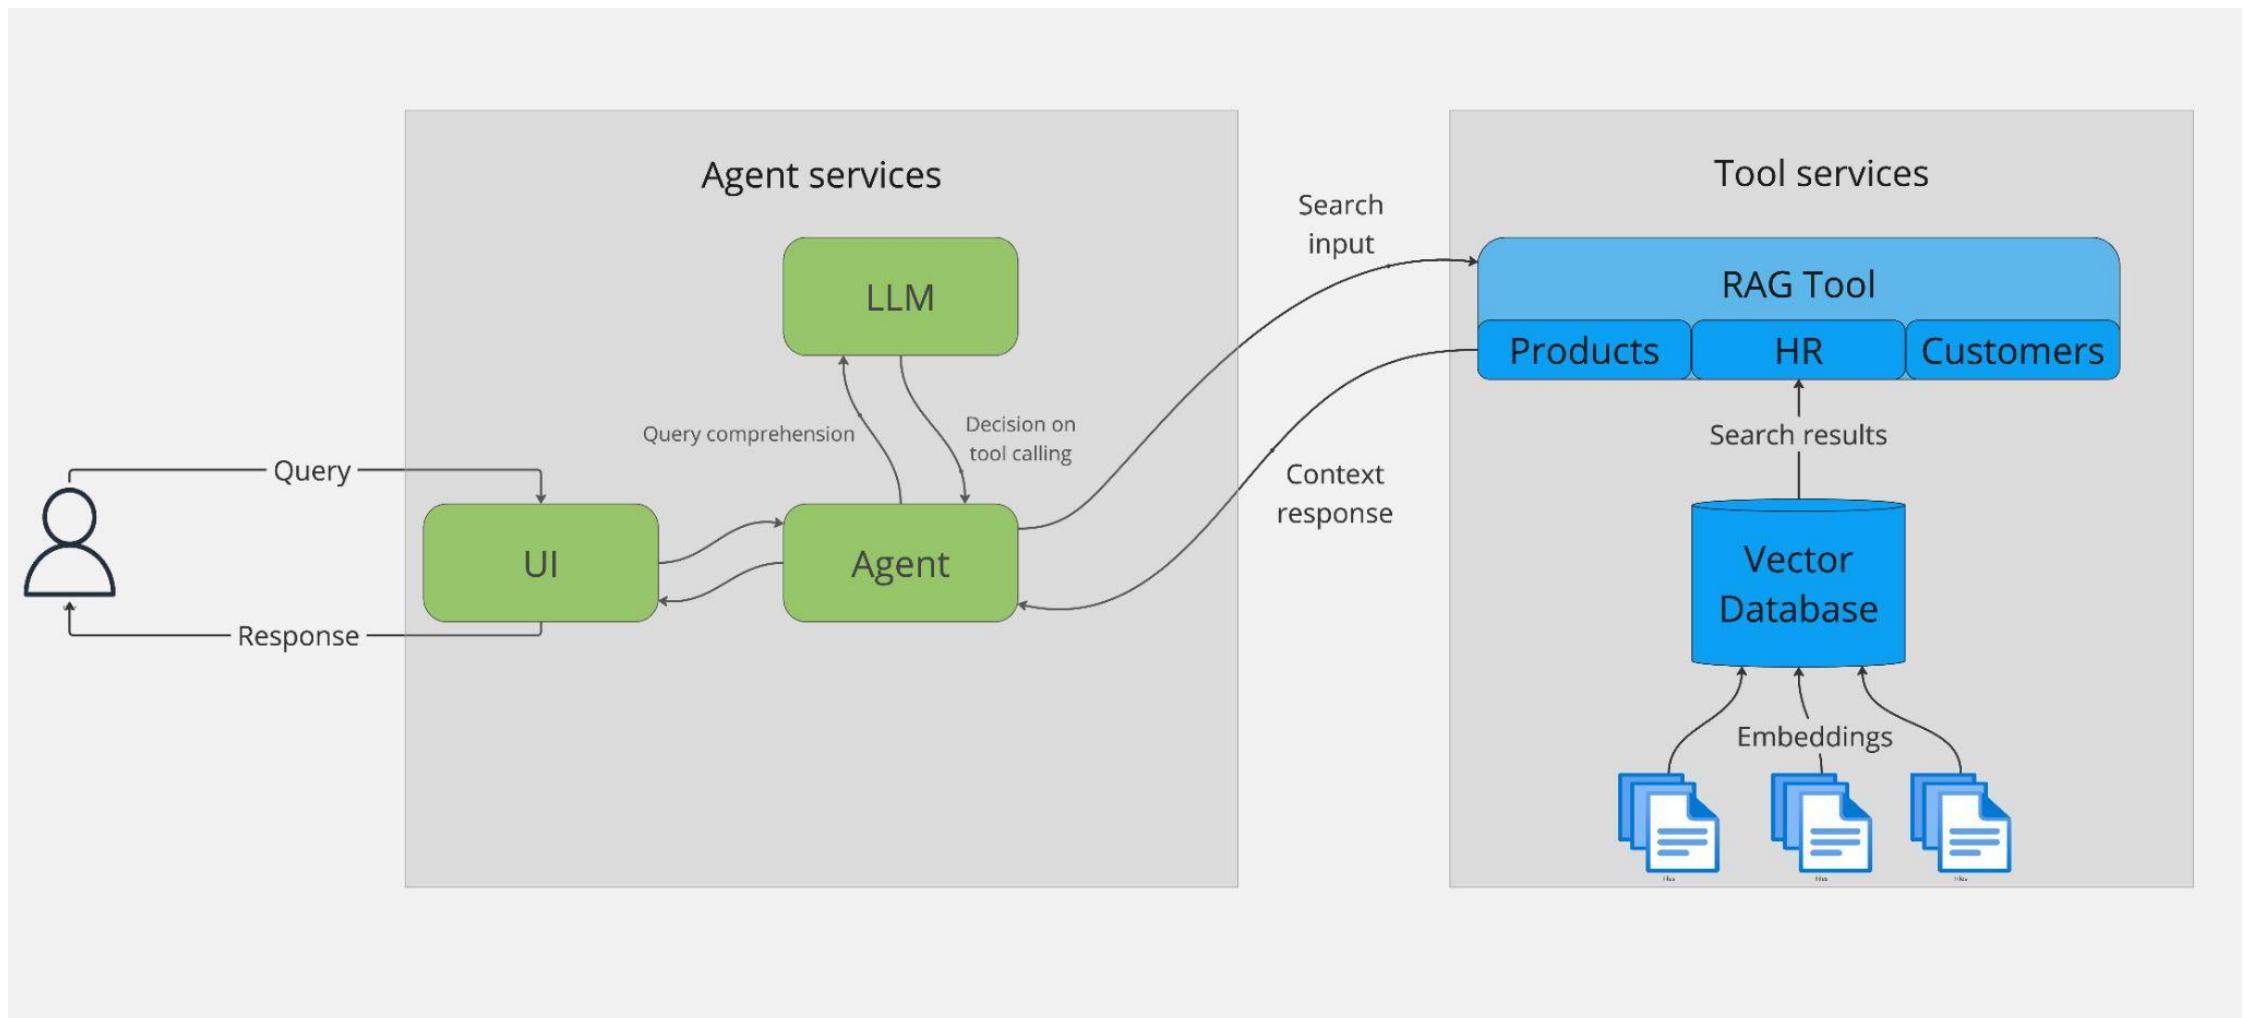
DEMO



Example Queries

- What does the CloudForge Migrate product do?
- What HR things should I do before the start date?
- Total Payments Received from FinNova Bank?

Architecture

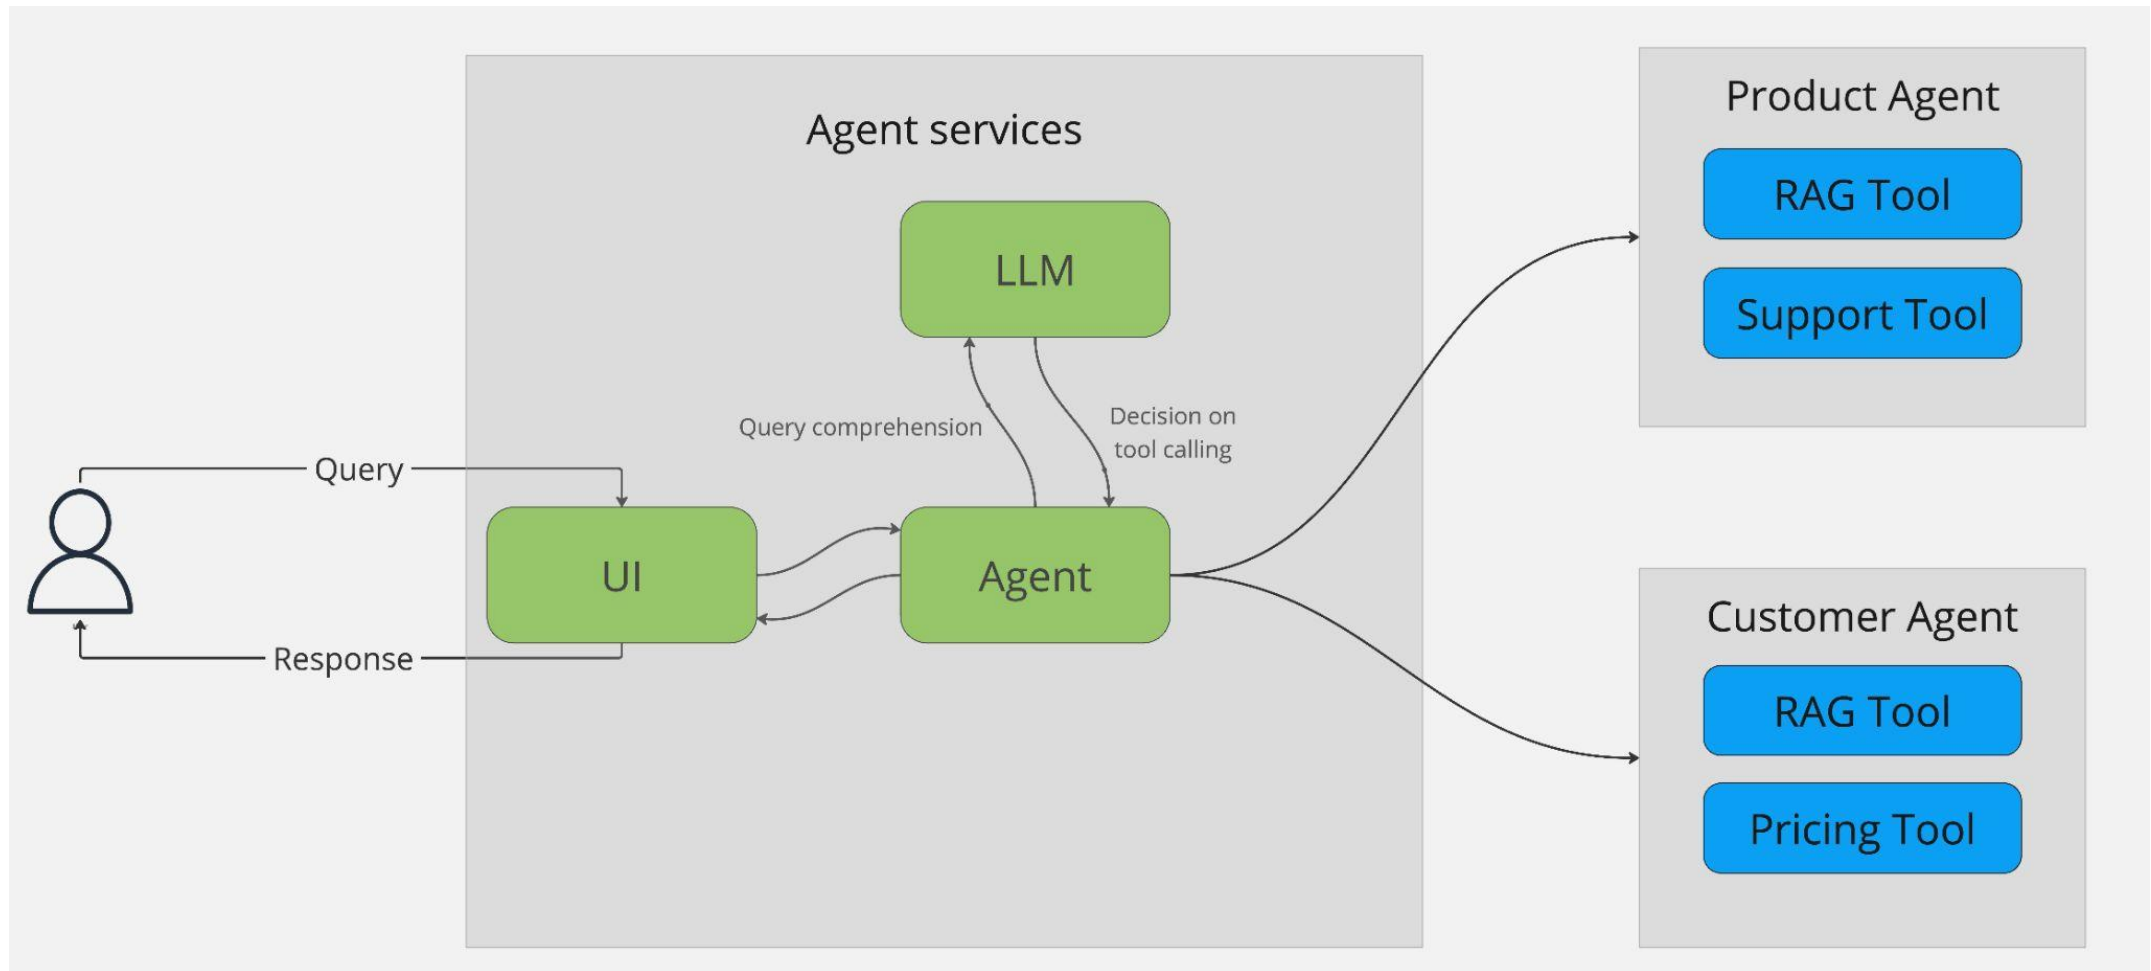


Resources

- **GitHub Repository** - <https://github.com/redhat-et/llm-agents>
- **Agent Services** (UI, Agent, LLM deployment files)
 - **Streamlit UI** - <https://github.com/redhat-et/llm-agents/tree/main/streamlit/openshift>
 - **ReAct Agent** - https://github.com/redhat-et/llm-agents/tree/main/react_agent/openshift
- **RAG Tool Services**
 - **Milvus vector database** - https://github.com/redhat-et/llm-agents/blob/main/rag/vector_db
 - **RAG service** - <https://github.com/redhat-et/llm-agents/blob/main/rag>



Discussion





Concluding Remarks

- LLMs are awesome, but they need RAG for grounding
- LLMs + RAG pattern is awesome, but it needs Agent's framework for planning and interacting with external sources
- **Challenges with the agents approach:**
 - Debugging is difficult as there can be many fault points such as selection of the right tool and input parameters
 - Latency of response is high if there are many API calls involved in creating the response
 - Load balancing and security concerns for external tools



KubeCon



CloudNativeCon

North America 2024



Thank You!

Q?



hveeradh@redhat.com
shanand@redhat.com